

Matching Continued

INFO/STSCI/ILRST 3900: Causal Inference

17 Oct 2024

Learning goals for today

At the end of class, you will be able to:

1. Understand propensity score matching and coarsened exact matching
2. Use matching methods to estimate causal effects

Matching: so far

Goal: Sample Average Treatment Effect on the Treated

$$E(Y^{a=1} | A = 1) - E(Y^{a=0} | A = 1)$$

Potential Solution: Create a group of untreated individuals, \mathcal{M} , which have a **similar distribution of L** to the treated group

$$\frac{1}{n_m} \sum_{i \in \mathcal{M}} Y_i = \frac{1}{n_m} \sum_{i \in \mathcal{M}} Y_i^{a=0} \approx \frac{1}{n_t} \sum_{i: A_i=1} Y_i^{a=0} \approx E(Y^{a=0} | A = 1)$$

How:

- Find untreated unit(s) which are similar to each treated unit

Matching: so far

Goal: Sample Average Treatment Effect on the Treated

$$E(Y^{a=1} \mid A = 1) - E(Y^{a=0} \mid A = 1)$$


Potential Solution: Create a group of untreated individuals, \mathcal{M} , which have a **similar distribution of L** to the treated group

$$\frac{1}{n_m} \sum_{i \in \mathcal{M}} Y_i = \frac{1}{n_m} \sum_{i \in \mathcal{M}} Y_i^{a=0} \approx \frac{1}{n_t} \sum_{i: A_i=1} Y_i^{a=0} \approx E(Y^{a=0} \mid A = 1)$$


How:

- ▶ Find untreated unit(s) which are similar to each treated unit
- ▶ Define “similar”

Motivating example



Motivating example



- ▶ \vec{L} denotes the set of covariates we want to condition on
- ▶ \vec{L}_i denotes the covariates values for individual i

A common distance metric: Exact matching


A common distance metric: Exact matching

- Ideally, we find an exact match for each treated unit

$$d(i, j) = \begin{cases} 0 & \text{if } \vec{L}_i = \vec{L}_j \\ \infty & \text{if } \vec{L}_i \neq \vec{L}_j \end{cases}$$

Because of the curse of dimensionality, this often leads to **no matches at all**

What if \vec{L} is multivariate? We need a **distance metric**



- ▶ Manhattan distance: $d(i,j) = \sum_p |L_{pi} - L_{pj}|$
- ▶ Euclidean distance: $d(i,j) = \sqrt{\sum_p (L_{pi} - L_{pj})^2}$
- ▶ Mahalanobis distance: Rescales and accounts for correlation

A common distance metric: Coarsened exact matching¹

¹Iacus, S. M., King, G., & Porro, G. (2012). [Causal inference without balance checking: Coarsened exact matching](#). Political Analysis, 20(1), 1-24.

A common distance metric: Coarsened exact matching¹

- Define $\tilde{\vec{L}}$ to be a coarsened version of \vec{L}

¹Iacus, S. M., King, G., & Porro, G. (2012). [Causal inference without balance checking: Coarsened exact matching](#). Political Analysis, 20(1), 1-24.

A common distance metric: Coarsened exact matching¹

- ▶ Define $\tilde{\vec{L}}$ to be a coarsened version of \vec{L}
 - ▶ Example: Age 15-20, 20-25, 25-30, etc

¹Iacus, S. M., King, G., & Porro, G. (2012). [Causal inference without balance checking: Coarsened exact matching](#). Political Analysis, 20(1), 1-24.

A common distance metric: Coarsened exact matching¹

- ▶ Define \tilde{L} to be a coarsened version of \vec{L}
 - ▶ Example: Age 15-20, 20-25, 25-30, etc
- ▶ Match exactly on \tilde{L}

¹Iacus, S. M., King, G., & Porro, G. (2012). [Causal inference without balance checking: Coarsened exact matching](#). Political Analysis, 20(1), 1-24.

A common distance metric: Coarsened exact matching¹

- ▶ Define \tilde{L} to be a coarsened version of L
 - ▶ Example: Age 15-20, 20-25, 25-30, etc
- ▶ Match exactly on \tilde{L}

$$d(i, j) = \begin{cases} 0 & \text{if } \tilde{L}_i = \tilde{L}_j \\ \infty & \text{if } \tilde{L}_i \neq \tilde{L}_j \end{cases}$$

¹Iacus, S. M., King, G., & Porro, G. (2012). [Causal inference without balance checking: Coarsened exact matching](#). Political Analysis, 20(1), 1-24.

A common distance metric: Coarsened exact matching¹

- ▶ Define \tilde{L} to be a coarsened version of L
 - ▶ Example: Age 15-20, 20-25, 25-30, etc
- ▶ Match exactly on \tilde{L}

$$d(i, j) = \begin{cases} 0 & \text{if } \tilde{L}_i = \tilde{L}_j \\ \infty & \text{if } \tilde{L}_i \neq \tilde{L}_j \end{cases}$$

- ▶ Benefit: Very transparent

¹Iacus, S. M., King, G., & Porro, G. (2012). [Causal inference without balance checking: Coarsened exact matching](#). Political Analysis, 20(1), 1-24.

A common distance metric: Coarsened exact matching¹

- ▶ Define \tilde{L} to be a coarsened version of L
 - ▶ Example: Age 15-20, 20-25, 25-30, etc
- ▶ Match exactly on \tilde{L}

$$d(i, j) = \begin{cases} 0 & \text{if } \tilde{L}_i = \tilde{L}_j \\ \infty & \text{if } \tilde{L}_i \neq \tilde{L}_j \end{cases}$$

- ▶ Benefit: Very transparent
- ▶ Benefit: Directly targets balance in L

¹Iacus, S. M., King, G., & Porro, G. (2012). [Causal inference without balance checking: Coarsened exact matching](#). Political Analysis, 20(1), 1-24.

A common distance metric: Coarsened exact matching¹

- ▶ Define \tilde{L} to be a coarsened version of L
 - ▶ Example: Age 15-20, 20-25, 25-30, etc
- ▶ Match exactly on \tilde{L}

$$d(i, j) = \begin{cases} 0 & \text{if } \tilde{L}_i = \tilde{L}_j \\ \infty & \text{if } \tilde{L}_i \neq \tilde{L}_j \end{cases}$$

- ▶ Benefit: Very transparent
- ▶ Benefit: Directly targets balance in L
- ▶ Drawback: May not find a good match for all individuals

¹Iacus, S. M., King, G., & Porro, G. (2012). [Causal inference without balance checking: Coarsened exact matching](#). Political Analysis, 20(1), 1-24.

A common distance metric: Coarsened exact matching¹


- ▶ Define \tilde{L} to be a coarsened version of L
 - ▶ Example: Age 15-20, 20-25, 25-30, etc
- ▶ Match exactly on \tilde{L}

$$d(i, j) = \begin{cases} 0 & \text{if } \tilde{L}_i = \tilde{L}_j \\ \infty & \text{if } \tilde{L}_i \neq \tilde{L}_j \end{cases}$$


- ▶ Benefit: Very transparent
- ▶ Benefit: Directly targets balance in L
- ▶ Drawback: May not find a good match for all individuals

¹Iacus, S. M., King, G., & Porro, G. (2012). [Causal inference without balance checking: Coarsened exact matching](#). Political Analysis, 20(1), 1-24.

A common distance metric: Propensity scores




A common distance metric: Propensity scores



Suppose \vec{L} only affects A through a probability of treatment

$$\pi_i = \pi(\vec{\ell}_i) = P(A_i = 1 \mid \vec{L} = \vec{\ell})$$


A common distance metric: Propensity scores



Suppose \vec{L} only affects A through a probability of treatment

$$\pi_i = \pi(\vec{\ell}_i) = P(A_i = 1 \mid \vec{L} = \vec{\ell})$$

A common distance metric: Propensity scores



Suppose \vec{L} only affects A through a probability of treatment

$$\pi_i = \pi(\vec{\ell}_i) = P(A_i = 1 \mid \vec{L} = \vec{\ell})$$

Conditional exchangeability holds given $\pi(\ell_i)$

Why propensity scores are nice

Why propensity scores are nice

- ▶ Can match on propensity scores directly instead of L

Why propensity scores are nice

- ▶ Can match on propensity scores directly instead of L
 - ▶ Easy to reason about

Why propensity scores are nice

- ▶ Can match on propensity scores directly instead of L
 - ▶ Easy to reason about
 - ▶ Can directly visualize the univariate matches

Why propensity scores are nice

- ▶ Can match on propensity scores directly instead of L
 - ▶ Easy to reason about
 - ▶ Can directly visualize the univariate matches
- ▶ Intuitive: Prioritizes covariates that predict treatment
- ▶ Mathematical guarantees on average

Why propensity scores are nice

- ▶ Can match on propensity scores directly instead of L
 - ▶ Easy to reason about
 - ▶ Can directly visualize the univariate matches
- ▶ Intuitive: Prioritizes covariates that predict treatment
- ▶ Mathematical guarantees on average
 - ▶ If our DAG is correct

Why propensity scores are nice

- ▶ Can match on propensity scores directly instead of L
 - ▶ Easy to reason about
 - ▶ Can directly visualize the univariate matches
- ▶ Intuitive: Prioritizes covariates that predict treatment
- ▶ Mathematical guarantees on average
 - ▶ If our DAG is correct
 - ▶ If our matches are good

Why propensity scores are nice

- ▶ Can match on propensity scores directly instead of L
 - ▶ Easy to reason about
 - ▶ Can directly visualize the univariate matches
- ▶ Intuitive: Prioritizes covariates that predict treatment
- ▶ Mathematical guarantees on average
 - ▶ If our DAG is correct
 - ▶ If our matches are good
 - ▶ We should **on average** get a matched group which looks like the the treatment group

$$P(L \mid \pi_i, A_i = 1) = P(L \mid \pi_i, A_i = 0)$$

Multivariate distances: Recap

When matching on multivariate \vec{L} , you have to define the distance between each pair of confounder values $\vec{\ell}_j$ and $\vec{\ell}_i$

- ▶ Manhattan distance
- ▶ Euclidean distance
- ▶ Mahalanobis distance
- ▶ Coarsened exact distance
- ▶ Propensity score distance

There is no right answer! Depends on the setting.

Multivariate distances: Recap

When matching on multivariate \vec{L} , you have to define the distance between each pair of confounder values $\vec{\ell}_j$ and $\vec{\ell}_i$

- ▶ Manhattan distance
- ▶ Euclidean distance
- ▶ Mahalanobis distance
- ▶ Coarsened exact distance
- ▶ Propensity score distance

There is no right answer! Depends on the setting.

- ▶ Propensity scores are most popular
- ▶ Sometimes they are substantively meaningful

Multivariate distances: Recap

When matching on multivariate \vec{L} , you have to define the distance between each pair of confounder values $\vec{\ell}_j$ and $\vec{\ell}_i$

- ▶ Manhattan distance
- ▶ Euclidean distance
- ▶ Mahalanobis distance
- ▶ Coarsened exact distance
- ▶ Propensity score distance

There is no right answer! Depends on the setting.

- ▶ Propensity scores are most popular
- ▶ Sometimes they are substantively meaningful
- ▶ Balance only occurs on average

Evaluate the matched sets

Whatever method, you should check that it worked

- ▶ Compare means of \vec{L} (propensity scores) across groups
- ▶ Possibly compare interaction cells; e.g., race \times age

Evaluate the matched sets

Whatever method, you should check that it worked

- ▶ Compare means of \vec{L} (propensity scores) across groups
- ▶ Possibly compare interaction cells; e.g., race \times age
- ▶ Visually assess distribution

Overlap

- ▶ Lack of overlap may indicate violation of positivity assumption

$$P(A = a \mid L = \ell) > 0 \text{ for all } a$$

- ▶ Ex: Sarah has no MD training. Would Sarah earn more money if she were a surgeon?

$$P(A = \text{Surgeon} \mid \text{No MD}) = 0$$

- ▶ If no good match exists, could be that $P(A = 0 \mid L = \ell) = 0$

Matching: A word of warning²

²Sekhon, J. S. (2009). [Opiates for the matches: Matching methods for causal inference](#). Annual Review of Political Science, 12(1), 487-508.

Matching: A word of warning²

$$L \rightarrow A \rightarrow Y$$


Matching: A word of warning²

Matching works!




Matching: A word of warning²

Matching works!



No help!




Matching: A word of warning²


Matching works!



No help!



No help!




²Sekhon, J. S. (2009). [Opiates for the matches: Matching methods for causal inference](#). Annual Review of Political Science, 12(1), 487-508.

Matching: A word of warning²


Matching works!



No help!



No help!



Matching is an estimation strategy.


²Sekhon, J. S. (2009). [Opiates for the matches: Matching methods for causal inference](#). Annual Review of Political Science, 12(1), 487-508.

Matching: A word of warning²


Matching works!



No help!



No help!



Matching is an estimation strategy.
It does not solve identification problems.

²Sekhon, J. S. (2009). [Opiates for the matches: Matching methods for causal inference](#). Annual Review of Political Science, 12(1), 487-508.

Estimating a causal effect

- ▶ If we've matched everything well, we can compare the means

³On the statistical role of inexact matching in observational studies. Guo and Rothenhäusler (2023)

Estimating a causal effect

- ▶ If we've matched everything well, we can compare the means
 - ▶ Treated group (with a match)
 - ▶ Matched control group

³On the statistical role of inexact matching in observational studies. Guo and Rothenhäusler (2023)

Estimating a causal effect

- ▶ If we've matched everything well, we can compare the means
 - ▶ Treated group (with a match)
 - ▶ Matched control group
- ▶ We can be extra careful by combining regression + matching
 - ▶ If everything is perfect, both should be fine on their own

³On the statistical role of inexact matching in observational studies. Guo and Rothenhäusler (2023)

Estimating a causal effect

- ▶ If we've matched everything well, we can compare the means
 - ▶ Treated group (with a match)
 - ▶ Matched control group
- ▶ We can be extra careful by combining regression + matching
 - ▶ If everything is perfect, both should be fine on their own
 - ▶ Combining can reduce bias
 - ▶ Reduces model sensitivity³

³On the statistical role of inexact matching in observational studies. Guo and Rothenhäusler (2023)

Learning goals for today

At the end of class, you will be able to:

1. Understand propensity score matching and coarsened exact matching
2. Use matching methods to estimate causal effects